

```
In [24]: import pandas as pd
import pickle
import warnings
warnings.filterwarnings("ignore")
```

```
In [25]: a = pd.read_csv("Health_insurance (1).csv")
```

```
In [26]: a
```

```
Out[26]:
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520
...
1333	50	male	30.970	3	no	northwest	10600.54830
1334	18	female	31.920	0	no	northeast	2205.98080
1335	18	female	36.850	0	no	southeast	1629.83350
1336	21	female	25.800	0	no	southwest	2007.94500
1337	61	female	29.070	0	yes	northwest	29141.36030

1338 rows × 7 columns

```
In [27]: a.describe()
```

```
Out[27]:
```

	age	bmi	children	charges
count	1338.000000	1338.000000	1338.000000	1338.000000
mean	39.207025	30.663397	1.094918	13270.422265
std	14.049960	6.098187	1.205493	12110.011237
min	18.000000	15.960000	0.000000	1121.873900
25%	27.000000	26.296250	0.000000	4740.287150
50%	39.000000	30.400000	1.000000	9382.033000
75%	51.000000	34.693750	2.000000	16639.912515
max	64.000000	53.130000	5.000000	63770.428010

```
In [28]: a.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 1338 entries, 0 to 1337  
Data columns (total 7 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   age         1338 non-null   int64  
1   sex         1338 non-null   object  
2   bmi         1338 non-null   float64  
3   children    1338 non-null   int64  
4   smoker      1338 non-null   object  
5   region      1338 non-null   object  
6   charges     1338 non-null   float64  
dtypes: float64(2), int64(2), object(3)  
memory usage: 73.3+ KB
```

```
In [29]: a.shape
```

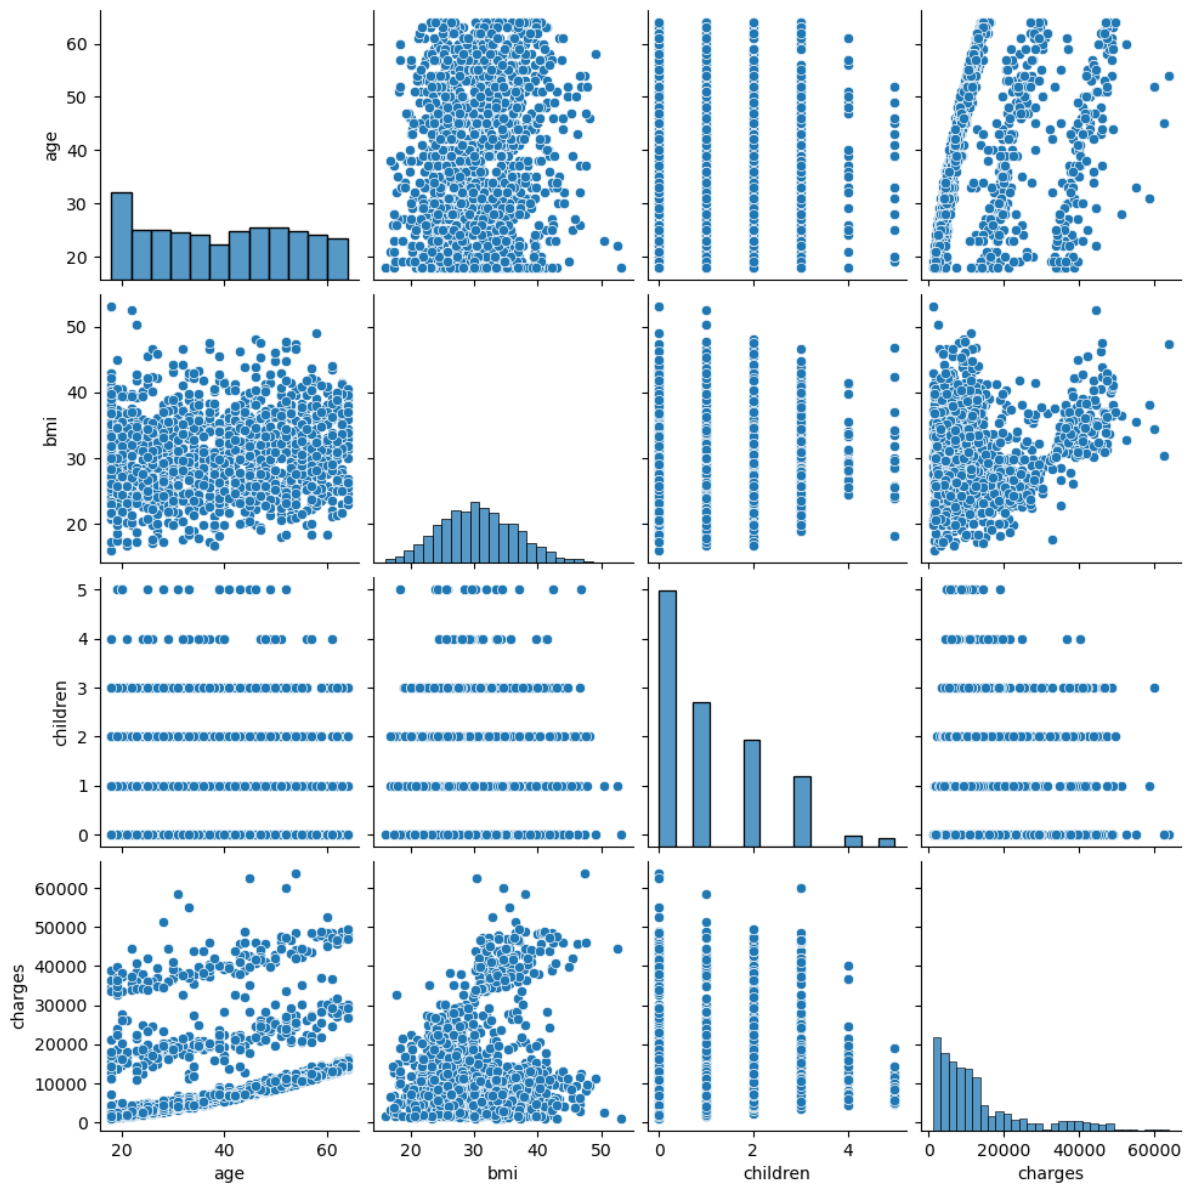
```
Out[29]: (1338, 7)
```

```
In [30]: a.isnull().sum()
```

```
Out[30]: age         0  
sex         0  
bmi         0  
children    0  
smoker      0  
region      0  
charges     0  
dtype: int64
```

```
In [31]: import seaborn as sns  
import matplotlib.pyplot as plt
```

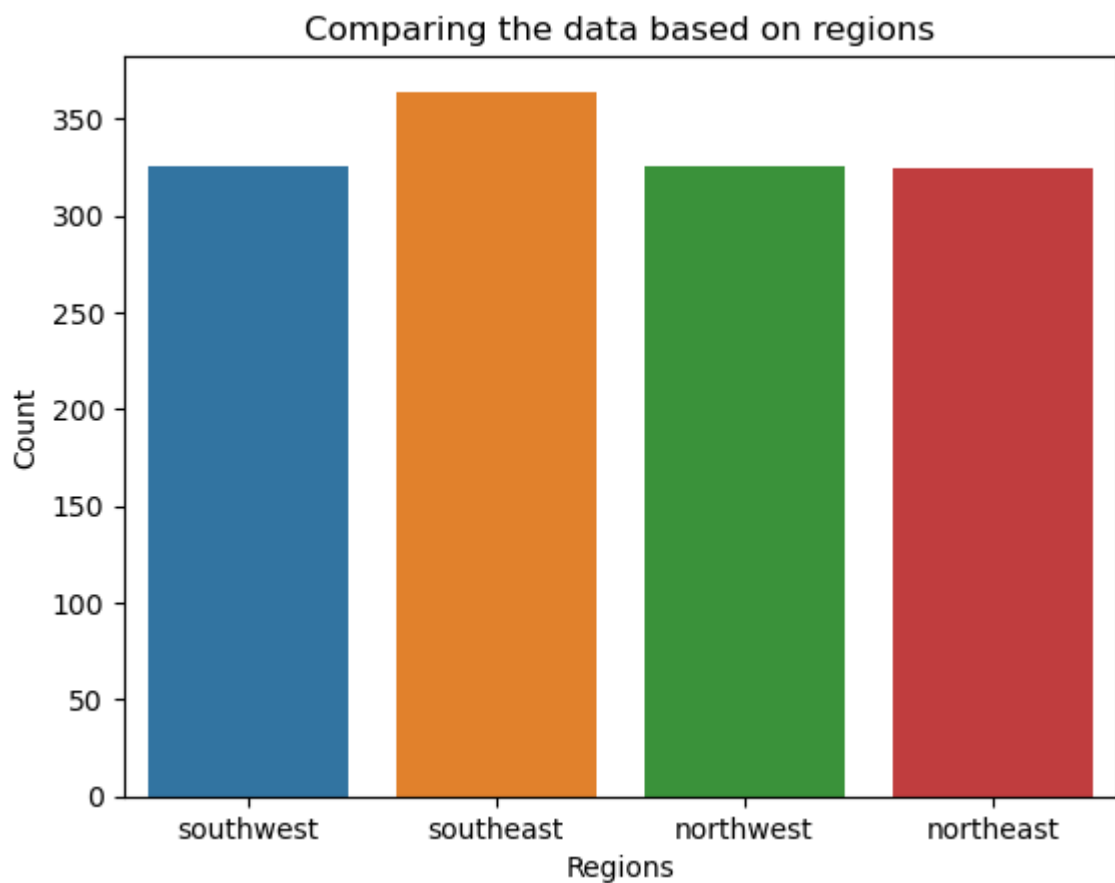
```
In [59]: sns.pairplot(a)  
plt.show()
```



```
In [61]: a["region"].value_counts()
```

```
Out[61]: southeast    364
southwest    325
northwest    325
northeast    324
Name: region, dtype: int64
```

```
In [62]: sns.countplot(x = a["region"])
plt.title("Comparing the data based on regions")
plt.xlabel("Regions")
plt.ylabel("Count")
plt.show()
```



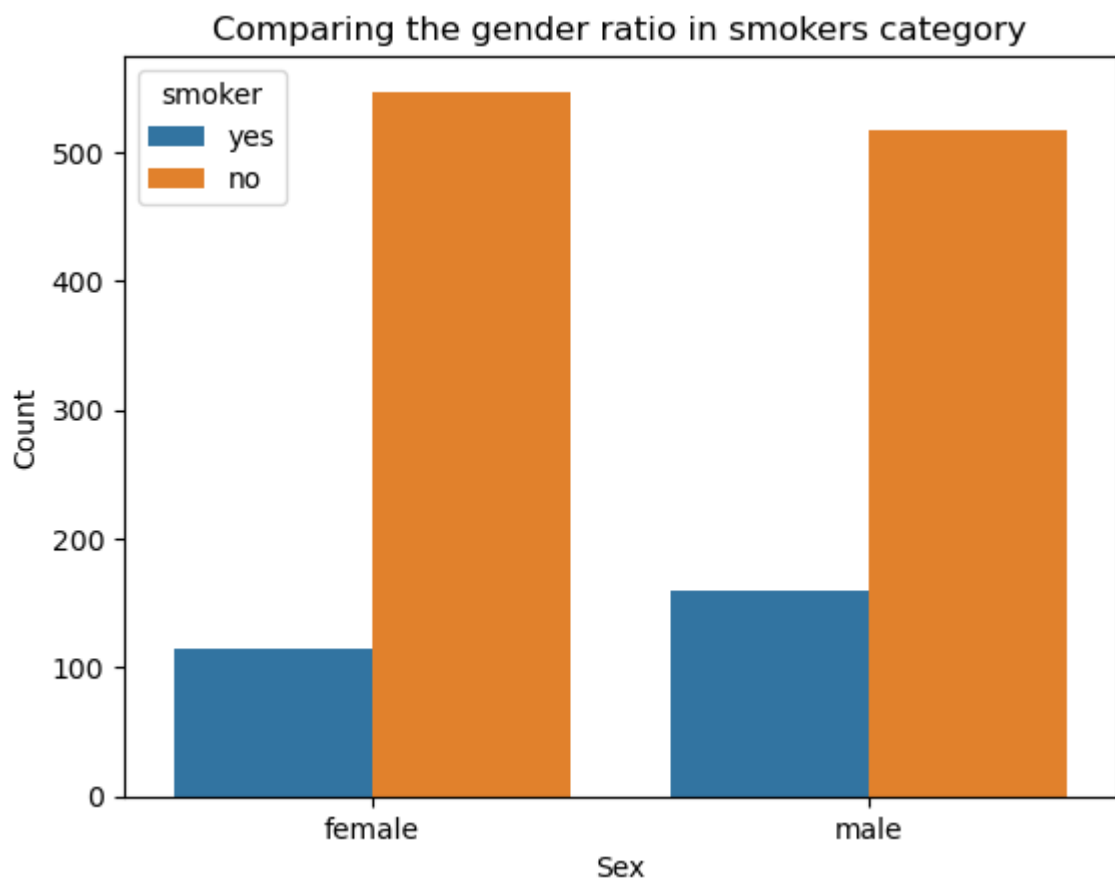
```
In [63]: a["sex"].value_counts()
```

```
Out[63]: male      676  
female    662  
Name: sex, dtype: int64
```

```
In [65]: a["smoker"].value_counts()
```

```
Out[65]: no      1064  
yes       274  
Name: smoker, dtype: int64
```

```
In [67]: sns.countplot(data=a, x="sex", hue="smoker")  
plt.title("Comparing the gender ratio in smokers category")  
plt.xlabel("Sex")  
plt.ylabel("Count")  
plt.show()
```



```
In [32]: data=pd.get_dummies(a,dtype=int)
```

```
In [33]: data.shape
```

```
Out[33]: (1338, 12)
```

```
In [34]: data.head(10)
```

```
Out[34]:
```

	age	bmi	children	charges	sex_female	sex_male	smoker_no	smoker_yes	region_north
0	19	27.900	0	16884.92400	1	0	0	1	
1	18	33.770	1	1725.55230	0	1	1	0	
2	28	33.000	3	4449.46200	0	1	1	0	
3	33	22.705	0	21984.47061	0	1	1	0	
4	32	28.880	0	3866.85520	0	1	1	0	
5	31	25.740	0	3756.62160	1	0	1	0	
6	46	33.440	1	8240.58960	1	0	1	0	
7	37	27.740	3	7281.50560	1	0	1	0	
8	37	29.830	2	6406.41070	0	1	1	0	
9	60	25.840	0	28923.13692	1	0	1	0	

```
In [35]: y=data['charges']
          x=data.drop('charges',axis=1)
```

```
In [47]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.25,random_state=11)
```

```
In [48]: x_test.head(5)
```

```
Out[48]:
```

	age	bmi	children	sex_female	sex_male	smoker_no	smoker_yes	region_northeast	region_south
1313	19	34.700	2	1	0	0	1	0	0
1254	34	27.720	0	1	0	1	0	0	0
372	42	33.155	1	1	0	1	0	1	1
937	39	24.225	5	1	0	1	0	0	0
484	48	34.300	3	0	1	1	0	0	0

```
In [49]: from sklearn.model_selection import GridSearchCV #GridSearchCV is for parameter tuning
from sklearn.ensemble import RandomForestRegressor
reg=RandomForestRegressor()
n_estimators=[25,50,75,100,125,150,175,200] #number of decision trees in the forest
criterion=['mse'] #criteria for choosing nodes default = 'gini'
max_depth=[3,5,10] #maximum number of nodes in a tree default = None (it will go to max_depth)
parameters={'n_estimators': n_estimators,'criterion':criterion,'max_depth':max_depth}
RFC_reg = GridSearchCV(reg, parameters)
RFC_reg.fit(x_train,y_train)
```

```
Out[49]: GridSearchCV(estimator=RandomForestRegressor(),
                      param_grid={'criterion': ['mse'], 'max_depth': [3, 5, 10],
                                   'n_estimators': [25, 50, 75, 100, 125, 150, 175, 200]}))
```

```
In [50]: RFC_reg.best_params_
```

```
Out[50]: {'criterion': 'mse', 'max_depth': 5, 'n_estimators': 175}
```

```
In [51]: reg=RandomForestRegressor(n_estimators=125,criterion='mse',max_depth=5)
```

```
In [52]: reg.fit(x_train,y_train)
```

```
Out[52]: RandomForestRegressor(criterion='mse', max_depth=5, n_estimators=125)
```

```
In [53]: y_pred=reg.predict(x_test)
```

```
In [54]: y_pred
```

```

Out[54]: array([36313.49832168, 6051.51233791, 7900.41727041, 7937.67247513,
10412.3678673 , 12403.6283807 , 6277.64195999, 14343.30322443,
14882.33840172, 44249.81917802, 4062.77278782, 13343.73081329,
12041.84455988, 2834.12461468, 7188.28390092, 25427.98836945,
4006.13893588, 2557.09887707, 7021.9161639 , 14106.49905248,
9365.56219324, 6956.19687256, 2500.35866915, 4308.2448154 ,
3699.66037698, 9807.96475179, 6563.13933865, 7254.44562817,
18699.83738138, 9616.43020131, 40362.21595493, 21291.20679336,
2664.59462964, 2631.64384729, 21826.95848098, 12335.74248548,
6013.82681346, 16836.40668246, 10862.09855014, 6569.68611999,
6941.44268142, 12627.07818441, 35122.45467305, 2393.603264 ,
9669.9189261 , 11921.23517353, 21303.61754544, 2500.35866915,
10048.01664037, 35458.47065285, 12729.08685913, 38506.68985557,
17324.17466636, 46436.35062487, 10472.89902437, 2355.73613423,
6055.05784912, 14370.01156659, 12463.8219875 , 8491.33303052,
7198.14714857, 13512.28659585, 12856.17994674, 48735.5470598 ,
8356.98657032, 7616.54453884, 14957.46307138, 35043.52854693,
9678.7539448 , 4480.71448355, 39885.12979704, 7609.31303965,
6951.33418237, 7109.10803502, 10177.16933216, 8960.19490078,
13539.99710094, 10648.13644681, 5912.52432664, 8539.55635905,
3662.22075512, 18736.7938785 , 5848.16302611, 41239.06282292,
8068.64998466, 8475.02989634, 7482.59603125, 3908.20801452,
23420.85705297, 13727.93010504, 14191.05343656, 4368.79741178,
7559.70699243, 14322.10917806, 7872.53208335, 13958.54192606,
4359.36222611, 17866.0959465 , 11981.71261561, 10752.12656594,
12046.84449792, 14291.35132816, 13104.60756802, 2526.91411074,
6067.51735657, 9635.70602386, 7583.56686625, 7712.34876029,
7208.50409621, 4400.34774558, 6830.06902848, 11479.02344809,
12321.16337157, 5788.98687665, 42395.37701945, 7306.00501997,
6037.10888844, 43010.11621789, 7144.92769436, 8228.22818152,
10061.04056519, 2363.41849767, 35439.37452351, 10543.15217227,
39643.82270666, 6954.70705226, 20275.24573511, 44381.61604454,
2819.40568323, 2673.5143155 , 35305.67310557, 9105.11012605,
6740.56333892, 8378.97944397, 14024.38931533, 7122.03549581,
4229.80677525, 7469.86316975, 17563.33810842, 2670.62900161,
7493.83327319, 13092.87466267, 11981.71261561, 12435.77394438,
8695.09231729, 6689.05282984, 10226.75584409, 46494.25914675,
13714.39271181, 44253.56976325, 21291.20679336, 6403.05666349,
2652.45514599, 6941.44268142, 14179.72637045, 43207.12525487,
7288.70964222, 14675.26248191, 4152.06444018, 12447.93393894,
23350.88219663, 2804.79949364, 7150.4505041 , 18451.68131695,
2673.5143155 , 7573.76732059, 18125.02918937, 27245.76513473,
9065.0062455 , 5851.96862414, 2646.25003687, 36298.71713636,
6489.27212974, 6901.16941935, 18801.42475425, 10913.4639419 ,
7171.83879692, 7495.73637003, 12677.50867809, 12142.28242813,
6157.15946239, 6695.85697842, 26329.02716395, 10506.8470944 ,
10445.94843506, 10659.43087603, 8885.84014538, 9661.10462036,
10513.10479488, 6600.07166598, 47182.47812368, 9218.04739214,
7648.93065515, 6572.51909658, 7664.17002724, 10912.03363079,
35817.78274345, 7129.44591912, 11476.27740821, 42692.11032902,
47055.78913112, 18155.58210003, 44935.78520674, 11872.08379126,
2834.12461468, 3742.00031943, 2716.58686298, 26331.02282402,
6436.49842205, 4282.95861034, 13743.89434131, 16145.87108556,
6942.52928658, 9097.94666386, 9431.03779621, 11944.83377469,
3696.03148827, 12039.14241618, 5854.90454739, 5426.51872528,
10885.11587288, 35997.63095791, 44080.64017615, 11304.93933293,
7117.74028414, 35458.47065285, 12714.29698755, 6836.74074822,
13321.72748318, 26759.98384881, 4686.60930148, 7697.97629726,
13878.41436193, 11216.40202106, 12352.43956111, 17857.48762498,
11158.44229825, 4326.50831005, 10415.5573113 , 4025.86925492,
6509.96850124, 6437.39082645, 7609.0323122 , 46861.30648317,
41083.18229876, 4143.30402881, 11892.68589197, 2518.53252477,
6263.83087449, 12550.75987885, 20047.76218608, 9988.89955438,
27368.24133547, 14240.90584694, 6883.94933902, 3974.1451684 ,

```

```

7114.46424286, 39883.58859567, 3997.43885783, 6865.03206507,
7261.3050877 , 7283.84091537, 5548.38163647, 2804.79949364,
6536.32702383, 2528.15143316, 8606.94701796, 11211.13882994,
10471.27125583, 3874.48657635, 10990.98515772, 14718.30393437,
6297.5920192 , 4070.67463158, 5744.08882908, 14153.71501863,
12532.22919217, 48811.06744707, 14848.56231146, 11421.98614211,
9654.42071912, 2393.603264 , 9792.03723751, 6777.02790228,
12392.42583847, 6653.95105999, 37396.27786352, 5511.18225555,
6732.15864627, 18392.39615544, 15654.45079018, 7000.59659045,
12465.45502876, 3943.41779229, 12091.23443857, 8066.84792806,
5914.86876479, 3742.00031943, 3803.47827797, 4377.84573543,
45114.22451736, 10386.27412988, 13261.37814423, 17834.05673859,
8611.12758971, 26288.7142562 , 2393.603264 , 9158.62802602,
14183.01241607, 44129.15506875, 19322.15387849, 4112.64005241,
4290.25391287, 16152.64080824, 8418.88530117, 2468.26588413,
15694.50014103, 44960.0890115 , 22120.3149577 , 7212.40589157,
13865.17291704, 2631.64384729, 14100.40276341, 39931.95451331,
12080.36901226, 6736.18388538, 18360.20774276, 12732.89203948,
8475.61238728, 49835.47506513, 7637.76989474, 14329.22359134,
7722.89415992, 7060.96093608, 7067.67460515])

```

```
In [55]: from sklearn.metrics import r2_score
r2_score(y_test,y_pred)
```

```
Out[55]: 0.8804199439174997
```


```
In [56]: res= pd.DataFrame(columns=['Actual', 'Predicted'])
res['Actual']=y_test
res['Predicted']=y_pred
#Results['km']=X_test['km']
res=res.reset_index()
res['Id']=res.index
res.head(10)
```

```
Out[56]:
```

	index	Actual	Predicted	Id
0	1313	36397.57600	36313.498322	0
1	1254	4415.15880	6051.512338	1
2	372	7639.41745	7900.417270	2
3	937	8965.79575	7937.672475	3
4	484	9563.02900	10412.367867	4
5	447	11454.02150	12403.628381	5
6	220	5012.47100	6277.641960	6
7	246	12648.70340	14343.303224	7
8	920	13451.12200	14882.338402	8
9	1070	39871.70430	44249.819178	9

```
In [57]: sns.lineplot(x='Id',y='Actual',data=res.head(100))
sns.lineplot(x='Id',y='Predicted',data=res.head(100))
plt.plot()
```

```
Out[57]: []
```

In []: